

In my inference engine I have divided the task of predicting if a person is wearing a mask and safety helmet into **two separate tasks**:

Task 1: First, an object detector network is used to predict bounding boxes containing heads in the image

Task 2: Next, two parallel image classifier networks are used to classify the detected head bounding boxes into 4 classes for the presence of mask(yes, no, invisible, wrong) and 2 classes for the presence of safety helmets(yes, no).

So, the **engine** consists of **two stages**:

Stage-1: Object detector network(Faster RCNN.) that takes an RGB image as input and produces bounding box locations for heads in the image. The object classifier branch of faster Rcnm classifies each ROI into **three classes** : **head, nonhead and background**. The object bounding box regressor branch produces 4 coordinates xtl, ytl, xbr, ybr specifying the bounding box location.

Stage-2 : 2 parallel Image classifier networks for mask and helmet.

Classifier 1(mask classifier) : Takes the bounding boxes labeled as head from the faster Rcnm detector and crops the bounding boxes from the image and takes them as input. The output has **4 classes(depending on values present in annotation) viz (invisible, no, wrong, yes)** denoting the presence of masks in the image. The architecture summary is shown below. Note that the final dense layer has 4 classes.

```
IPython: C:\prog\interview_data
(359, 224, 224, 3) (1433, 224, 224, 3)
(359, 4) (1433, 4)

Layer (type)                Output Shape                Param #
=====
input_7 (InputLayer)        (None, 224, 224, 3)        0
resnet50 (Model)            multiple                    23587712
global_average_pooling2d_1 ( (None, 2048)                0
dropout_1 (Dropout)         (None, 2048)                0
dense_5 (Dense)              (None, 5)                   10245
dropout_2 (Dropout)         (None, 5)                   0
dense_6 (Dense)              (None, 4)                   24
=====
Total params: 23,597,981
Trainable params: 23,544,861
Non-trainable params: 53,120

None
Train on 1003 samples, validate on 430 samples
Epoch 1/1
40/1003 [>.....] - ETA: 1:11:34 - loss: 1.4372 - acc: 0.2250
```

Mask classifier

Classifier 2(safety helmet classifier) : Takes the bounding boxes labeled as head from the faster Rcn and crops the bounding boxes from the image and takes them as input. The output has **2 classes(depending on values present in annotation) viz (yes, no)** denoting the presence of safety helmets in the image. The architecture summary is shown below. Note that the final dense layer has 2 classes.

IPython: C:\prog\interview_data

```
(1792, 224, 224, 3)
(359, 224, 224, 3) (1433, 224, 224, 3)
(359, 2) (1433, 2)
2020-06-04 05:33:33.621876: I T:\src\github\tensorflow\tensorflow\core\platform\cpu_feature_guard.cc:140] Your CPU supports instructions that this TensorFlow binary was not compiled to use: AVX2
```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	(None, 224, 224, 3)	0
resnet50 (Model)	multiple	23587712
global_average_pooling2d_1 ((None, 2048)		0
dropout_1 (Dropout)	(None, 2048)	0
dense_1 (Dense)	(None, 5)	10245
dropout_2 (Dropout)	(None, 5)	0
dense_2 (Dense)	(None, 2)	12

Total params: 23,597,969
 Trainable params: 23,544,849
 Non-trainable params: 53,120

None
 Train on 1003 samples, validate on 430 samples
 Epoch 1/1
 40/1003 [>.....] - ETA: 53:01 - loss: 0.7352 - acc: 0.4750

Helmet classifier